

Practical Depth Estimation with Image Segmentation and Serial U-Nets

Kyle J. Cantrell^a, Craig D. Miller^b and Carlos W. Morato


*Department of Robotics Engineering, Worcester Polytechnic Institute, 100 Institute Road, Worcester, MA, U.S.A.
{kjcantrell, cdmiller, cwmorato}@wpi.edu*


Keywords: Autonomous Vehicles, Depth Estimation, Ensemble Neural Networks, Intelligent Transport Systems, Semantic Segmentation, U-Net, Vehicle Perception, VSLAM.

Abstract: Knowledge of environmental depth is required for successful autonomous vehicle navigation and VSLAM. Current autonomous vehicles utilize range-finding solutions such as LIDAR, RADAR, and SONAR that suffer drawbacks in both cost and accuracy. Vision-based systems offer the promise of cost-effective, accurate, and passive depth estimation to compete with existing sensor technologies. Existing research has shown that it is possible to estimate depth from 2D monocular vision cameras using convolutional neural networks. Recent advances suggest that depth estimate accuracy can be improved when networks used for supplementary tasks such as semantic segmentation are incorporated into the network architecture. A novel Serial U-Net (NU-Net) architecture is introduced as a modular, ensembling technique for combining the learned features from N-many U-Nets into a single pixel-by-pixel output. Serial U-Nets are proposed to combine the benefits of semantic segmentation and transfer learning for improved depth estimation accuracy. The performance of Serial U-Net architectures are characterized by evaluation on the NYU Depth V2 benchmark dataset and by measuring depth inference times. Autonomous vehicle navigation can substantially benefit by leveraging the latest in depth estimation and deep learning.

1 INTRODUCTION

Typical color monovision cameras provide three data points per pixel in a given image. The data points correspond to the red, green, and blue (RGB) intensity levels (from 0 to 255) present in the image at that specific point. Stereo vision cameras are able to provide a fourth datapoint at each pixel, namely depth. Since depth is required for vehicles to localize and build maps of their environment, many autonomous vehicles rely on stereo vision. Although stereo vision cameras can provide the required depth data, they are frequently orders of magnitude more expensive than a single monovision camera. Achieving stereo vision levels of performance from a monovision camera is desirable as they are smaller, cheaper, and less sophisticated. As the cost of computational power continues to plummet, the additional processing power required to use a neural network to extract the depth data from a monovision camera feed does not present a large burden. This paper explores the estimation of depth data strictly from 2D RGB values from a monovision camera.

^a  <https://orcid.org/0000-0002-2785-9881>

^b  <https://orcid.org/0000-0003-2562-6088>




Figure 1: Depth prediction results on KITTI 15 (Godard et al., 2017).

1.1 Problem Description

An important problem in intelligent transport systems is autonomous navigation, an inevitably complex super-set of problems involving object detection, trajectory generation, and the colloquially termed ODOA (Obstacle Detection and Object Avoidance). At the heart of this problem is the ability to accurately estimate the depth of an object within a field of view (FOV). This is achievable today with several

techniques such as stereo-vision (using the overlap of two cameras in a computation to triangulate the position of an object in a frame relative to the camera resolutions and position). Another method is pairing a camera with a range finding solution like LIDAR, RADAR or ultrasonic sensors. Both of these methods have drawbacks in hardware cost and computational cost.

In addition to producing reliable depth data is the need to produce fast, actionable depth data as well. Without sufficiently fast depth data processing, the decision systems in place will be unable to safely navigate without the aid of an online, collaborative environment which is unrealistic for practical autonomous vehicles. The true nature of autonomy in regards to navigation is the independent decision making that can take place and in order for that to happen in a safe, robust way, timely and accurate data is needed in a continuous stream format.

The purpose of this paper is to present a solution to the problem of ODOA via depth estimation with monocular devices, thereby reducing the amount of hardware and computational cost necessary for an autonomously navigated robot.

1.2 Literature Review

Many researchers have attempted to solve this problem with various levels of success. Early naive approaches involved hand-crafted features and resulted in only modest accuracy. More recent approaches have trained convolutional neural networks (CNNs) to simultaneously predict depth and offer semantic segmentation of an image. A few datasets exist for benchmarking depth estimation. The two most popular sets are KITTI (Geiger et al., 2013) and NYU Depth Dataset V2 (NYUD v2) (Silberman et al., 2012). Both of these references have been cited thousands of times and are the de facto standards for validating new depth estimation frameworks. State-of-the-art performance on the KITTI dataset is currently a relative square error of 2.00. This result was achieved by Manuel López Antequera (Mapillary, September 2019) and is not yet published. Current state-of-the-art performance on the NYUD v2 dataset is a root-mean-square error of 0.356 (Lee, 2019). Performance on both datasets has rapidly improved over the past several years due to the experimentation and implementation of a variety of techniques. Traditional depth estimation networks take RGB values from stereo images as input and the depth values as the ground truth for backpropagation. Once the network is trained, the network can be presented monocular images and it is able to estimate depth.

The networks typically utilize an end-to-end, pixel-to-pixel architecture. This means that the size of the input and output vectors will be identical. In (Jiao et al., 2018) and (Wang et al., 2015), semantic segmentation was shown to improve depth estimation accuracy. (Jiao et al., 2018) uses a hybrid network architecture and defines an attention-driven loss to improve efficacy. Semantic segmentation is a rich field of study on its own and can be approached in many ways. Even though it is a field in its own right, semantic segmentation is an integral part to our objective as the semantic scene labeling is important to our vision of contributing to solving the obstacle detection and avoidance problem. Using methods presented in (Farabet et al., 2012), we hope to utilize semantic segmentation to inform our convolutional network on the correct depth estimation predictions.




Figure 2: Original RGB Image and Corresponding Segmentation.




Figure 3: Network Architecture Combining Depth Estimation and Semantic Segmentation (Jiao et al., 2018).

Given the basis of existing research, we hypothesize that by combining many of the latest techniques, we will be able to achieve near state-of-the-art depth estimation accuracy performance on NYUD v2. Specifically, this model will:

- Leverage the existing U-Net architecture
- Pre-train encoder layers on the ImageNet dataset
- Incorporate state-of-the-art semantic segmentation
- Allow for the integration of networks trained on supplementary tasks

The network will be benchmarked on the NYUD v2 dataset. If the network produces accurate output, it could serve as a basis for ODOA and 3D mapping algorithms simply using a monocular camera, the impli-

cations of which are current technologies that enable ODOA, such as LIDAR, RADAR and Sonar, could be enhanced or replaced. LIDAR, in particular, has many drawbacks but finds itself as one of the more prevalent technologies of choice for this application because of its price point. Some of the drawbacks to this can be seen in (Farabet et al., 2012).

Current systematic solutions to depth estimation tend to fall short in several areas including accuracy, scalability and Size, Weight, Power and Cost (SWaPC). For instance, the current state-of-the-art performance holder on KITTI (Lee, 2019), has a relative squared error of 2.00, inferring a 98 % efficacy rate. At short distances and low velocities this isn't a problem in autonomous robotic applications but at high velocities in sensory heavy environments, 2% can be the difference of collision or mission success when factored in with inference time, processing time for other tasks, and state machine updates. Further, considering again the state-of-the-art holder, the base models used in the paper are based heavily on exceedingly deep architectures (ResNet50/100, DenseNet121, etc). The relationship between accuracy and inference times are proportional in that the higher the accuracy (and thus deeper the network) the higher the inference time. In the case of a base ResNet50, the shallowest architecture, takes 103ms for an error rate of 7 as shown in (He et al., 2016). The proposed solution seeks to bridge the accuracy, SWaPC, and speed divide by leveraging state-of-the-art deep learning to extend the ability of standard cameras into scene-understanding sensors with intelligent comprehension suitable for autonomous driving decision making. Based on the current research referenced in this section, it is evident that performing depth estimation through the use of artificial neural networks to decompose an image into its essential depth features can provide a depth output mapping suitable for near-real-time inferences without the cost associated with expensive LIDAR, RADAR, or Stereovision systems, with the added benefit of integrating with these systems in the future for combinatorial improvement.

2 EXISTING DEPTH ARCHITECTURES

Several existing neural network architectures are trained to predict depth from RGB images.

2.1 Standard CNN

Standard CNN architectures are a default choice for many computer vision tasks such as image classification, digit recognition, and facial recognition. For this application, a sequential CNN with two convolutional layers followed by alternating densely connected and dropout layers was evaluated. Nearly all convolutional and densely connected layers were activated by the ReLU nonlinearity. The output layer was linearly activated.

2.2 RCNN

Recurrent Convolutional Neural Networks (RCNN) represent a hybrid architecture that can leverage the state-of-the-art in both sequence-based deep learning (RNN) and image classification (CNN). Several commercial and high performance RCNN models have been successful in related fields such as the R-CNN-192 developed by Ming Liang et al in Recurrent Convolutional Neural Network for Object Recognition (Liang and Hu, 2015). The idea to combine the proven standard Convolutional Neural Network classifier and segmentation ability with a proven sequential predicting Recurrent network was born of the idea that the depth data being trained can inform the LSTM layer(s) that a desired feature to be learned beyond the object segmentation is the relation between the front and the back of the image. An example architecture can be seen in Figure 4.




Figure 4: CNN + LSTM Model.

For practical purposes, this is a standard Convolutional Neural Network with a single LSTM layer of 512 Units. The input to the neural network is a single RGB image with a resolution of 640x480 pixels. After two convolutional layers with two pooling layers interleaved in between, a 50 percent dropout layer is inserted. The output of the dropout layer is connected to the LSTM layer input with 1 sample, timestep of 1 and 512 features (the output of the CNN). An alternative RCNN was developed as well where a second hidden LSTM layer was constructed.

2.3 U-Net

U-Net is an encoder-decoder neural network architecture. The 2D size of U-Net's input and output arrays

are equal. U-Net is frequently used for the segmentation of images. Since depth estimation requires a 2D array output equal to the number of input pixels, U-Net is an appropriate model candidate. For segmentation, U-Net is typically appended with a final activation layer for assigning semantic classes. In depth estimation, the output of the final decoding convolutional layer can be taken as a 2D array of 8-bit depth values with which a depth image can be constructed.

U-Net gets its name from the structure of the net itself. The first half of the “U” represents the encoder and the second half represents the decoder. Each step on each side of the U represents a block of convolutional and pooling layers. The output of each of the encoding blocks feeds the next encoding block as well as the analogous decoding block. Higher and lower level image features are available at different blocks of the network by connecting the weights across the U.




Figure 5: U-Net Architecture with ResNet34 Backbone (Ronneberger et al., 2015).

U-Net architectures can be built using a number of different base models. ResNet, MobileNet, and VGG are popular base model choices. In this experiment, a U-Net architecture with a ResNet34 backbone was loaded with encoder weights pre-trained on the ImageNet dataset. The encoder weights were frozen during the depth estimation training to preserve the feature extractors learned for image segmentation.

All existing architectures offer some benefit to the depth estimation problem. Each of the models solve part of the problem, but not the entirety of it. For instance, Standard Convolutional Neural Networks perform admirably on object recognition and classification, but lack an awareness of varying depths with pixels. Recurrent Convolutional Neural Networks should address the time varying output problem but fall short in object detection. The best performing object detection models usually involve an encoder-decoder “U-Net” style architecture and these models too fail when trying to track the depth within images because of an impulsive training regimen. A novel ap-

proach to solving these drawbacks is an ensemble architecture combining the best parts of the existing architectures mentioned in this section into an arrangement termed “NU-Net”, a portmanteau including the number of N U-Nets connected serially. The development of this architecture is detailed in Section 3.

3 SERIAL U-NETS

A Serial U-Net, or NU-Net, is an ensemble network architecture for combining the learned features from N-many U-Nets into a single pixel-by-pixel output. The key benefit of this architecture is the ability to enhance the performance of an overall networks primary task by integrating component U-Nets that were pre-trained on supplementary tasks. For example, a simple serial U-net (2U-Net) may include a pre-trained U-Net for semantic segmentation and a second U-Net for depth estimation. The networks primary task (depth estimation) is enhanced by the addition of the pre-trained segmentation network. A more advanced example may include three networks (3U-Net), two of which were pre-trained - one for semantic segmentation and another for object detection. Again, in this case the functionality of the pre-trained component networks are used to improve performance of the ensemble networks primary task of depth estimation. In general, Serial U-Nets provide a modular architecture for integrating supplementary learned features into a new pixel-to-pixel ensemble network.

Two main variations of serial U-Nets are presented in this paper. The first variation, NU-Net, simply takes the output of the first component U-Net as input to the following component U-Net. When integrating pre-trained networks, it may be necessary to remove the final activation layer depending on the networks original task. In this way, the learned features are extracted from the component network and can be provided to the ensemble network in a meaningful way. Therefore, the “output” of a component U-Net in a Serial U-Net is not always guaranteed to be the exact output of its final layer as seen when operating as a standalone network.

The second variation presented is NU-Net Connected, which takes the output of the previous component U-Net concatenated with the original input image as input to the following component U-Net. NU-Net Connected seeks to eliminate any residual information loss from passing through the first component network. This is a critically important difference from the vanilla NU-Net. NU-Net can be viewed as a series of isolated functions that simply pipe outputs to each other, whereas NU-Net Connected is able to use

output from pre-trained component networks in addition to the original input image to perform its task. Reverting to the previous example, this implies that a 2U-Net Connected model could utilize the results of a component semantic segmentation network in addition to an unmodified RGB image to produce a depth estimate. This concept can be extended to include N-many U-Nets in series for the integration of additional learned features that can help produce more accurate predictions. Network sequence must be considered when designing both NU-Net and NU-Net Connected models.

Each component U-Net is structured based on a backbone architecture. The backbone architecture dictates the layer structure in the encoder half of the U-Net, which is then mirrored in the decoder half as well. Clearly, selection of each backbone architecture largely influences performance results. Any backbones can be used in the component U-Nets in a Serial U-Net. It is not necessary to utilize the same backbone for each component U-Net in a Serial U-Net model. Since the number of parameters in higher order Serial U-Nets can become large quickly, it is preferable to use component U-Nets with the smallest backbone that is able to achieve satisfactory results for that component task (segmentation, object detection, etc.). Attention must still be paid to matching the input and output layer sizes to ensure proper piping of each component result. Minimizing backbone sizes saves the level of computation required for inferencing. By using larger backbones, it is possible to increase the accuracy of the serial U-net at the expense of training time and inference time. When using larger backbones such as DenseNet201, there are many more parameters that must be calculated than in smaller backbones such as VGG16. Updating the additional parameters during training requires more computational resources results in longer training sessions.

In the Serial U-Net built for the task of depth estimation, the weights of the layers from the pre-trained component networks are frozen. The original functionality of the network is thus preserved. Without freezing the pre-trained weights, large gradients at the first few epochs of the depth estimation training process have the potential to dramatically change the initial weight values and possibly destroy the learned behavior. In this paper, weights in the final component U-Net module of the Serial U-Nets were previously untrained and randomly initialized. Pre-training and/or strategic weight initialization would likely result in improved performance.

To summarize, the Serial U-Net is a modular ensemble architecture for combining the learned fea-

tures of N-many component U-Nets. By combining component U-Nets pre-trained on supplementary tasks, performance on a primary task can be improved. Different backbone architectures can be selected for each component U-Net to attempt to optimize total computation time and ensemble performance on the primary task.

3.1 2U-Net (W-Net)

2U-Net, or W-Net, is a proposed Serial U-Net architecture for leveraging semantic segmentation for improved depth estimation. W-Net is shown in Figure 6. W-Net is composed of two U-Nets in series. It is the simplest possible implementation of a Serial U-Net. The first U-Net is a pre-trained network to perform segmentation. The output from the first U-Net is fed into the beginning of a second untrained U-Net. During the depth estimation training phase, the weights of all of the layers in the first U-Net are frozen.

3.2 2U-Net Connected (W-Net Connected)

W-Net Connected is also a proposed Serial U-Net architecture. W-Net Connected is shown in Figure 7. It is very similar to W-Net with the exception of a connection between the RGB input to the input of the second U-Net. In this way, the second U-Net sees the original image as well as the output from the first U-Net. By connecting the second U-Net to both of these layers, W-Net Connected is able to utilize a pre-segmented image (from U-Net #1 output) to help inform the depth estimate without losing any information from the directly connected original image.

After concatenating the RGB input with the first U-Net output, it is reshaped to (480,640,4) and passed to the second U-Net.

4 LOSS FUNCTIONS

During the research phase, it became apparent that experimenting with the Loss functions was necessary. As noted in (Eigen et al., 2014) and (Lee and Kim, 2019), Scale Invariant Error is a log-based objective function that works by penalizing in log steps the percent error of the predicted output versus the ground truth. In addition, it penalizes less when the direction of the output is consistent with the direction of the ground truth. The computation can be seen in (1).

$$L(y, y^*) = \frac{1}{n} \sum_i^n (d_i)^2 - \frac{\lambda}{n} \left(\sum_i^n (d_i) \right)^2 \quad (1)$$




Figure 6: Serial U-Net Architecture: 2U-Net (W-Net) with ResNet34 Component Backbone.

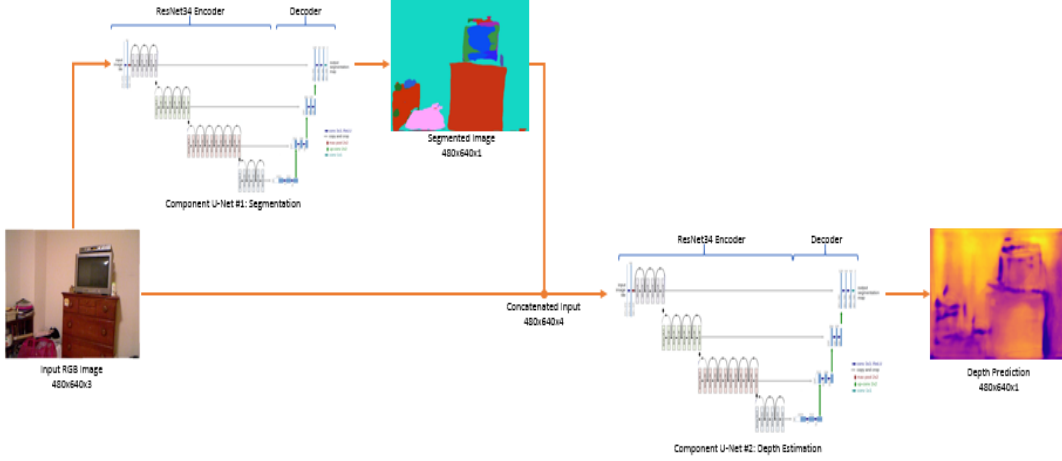


Figure 7: Serial U-Net Architecture: 2U-Net Connected (W-Net Connected) with ResNet34 Component Backbone.

Where y is the predicted output, y^* is the ground truth, $d_i = \log(y_i) - \log(y^*)$, and λ is a real number between $[0,1)$. Noteworthy for this algorithm is that the raw depth training data does not react well when trained with this loss function because of the grayscale and discrete nature of the image, even when normalized. To deal with this, the log of the predicted and ground truth pixels were “clipped”, or, limited to the value of ϵ ($1e-07$). It was determined during research that this loss function was not compatible with the training data in an unprocessed, or raw format.

Mean Squared Error was eventually chosen as the best performing Loss function and standard performance metric to measure. The equation for Mean Squared Error can be seen in (2).

$$L(y, y^*) = \frac{1}{n} \sum_i^n (y_i - y_i^*)^2 \quad (2)$$

This is a straightforward calculation measuring the error, per-pixel, between the ground truth and the predicted output image. It is worth noting that the state-of-the-art in this problem space is generally measured in Linear Root Mean Squared Error, Root-Mean-Squared-Log-Error, and Absolute Relative Tolerance.

5 METHODOLOGY

The existing U-Net and Serial U-Net architectures were benchmarked using the linear Mean Squared Error (MSE) loss function and the Adam Optimizer, an adaptive optimization algorithm. This is an appropriate selection for a loss function given that depth estimation is a regression task. Also evaluated was Scale Invariant Loss, a popular log error calculation that includes a “directional” term to correct the gradients with a finer granularity than mean squared error alone. This can be seen in depth in (Eigen et al., 2014). A learning rate of 0.0001 was used on all evaluations. All models were evaluated using a subset of the NYUD v2 dataset.

5.1 Evaluation Setup

All models were trained on 1,088 image pairs and validated on 360 image pairs from the NYUD v2 dataset. All training and test data - both RGB and depth - were normalized from a 0-255 to 0-1 scale. All models were trained on a NVIDIA RTX 2060.

Table 1: Regression Results, NYUD v2: 20 Epochs.

Regression Performance		
Architecture	MSE	RMSE
U-Net	0.0436	0.2087
W-Net	0.0421	0.2053
W-Net Connected	0.0412	0.2031

Table 2: Inference Benchmarks, NYUD v2: 2 Depth Epochs.

Inference Benchmarks (ms)			
Architecture	Average	Min.	Max.
U-Net	100.64	80.16	458.64
W-Net	133.02	120.15	360.60
W-Net Connected	132.66	118.58	195.70

6 RESULTS AND DISCUSSION

6.1 Performance Benchmarks

All architectures were evaluated by training them for 20 epochs. All networks were trained with a batch size of 2. A small batch size was used to prevent memory errors during training. Minimum validation MSEs are logged in the tables below from the 20 epoch training process.

After training, inference times for each model were measured on a desktop with an i5-4460 processor, 16 GB of RAM, and a NVIDIA RTX 2060. Average, minimum, and maximum inference times were logged after performing 100 model predictions. All models used the same RGB input image for each iteration.




Figure 8: Training and Validation Loss.



Figure 9: Input RGB Image for W-Net Connected Architecture.




Figure 10: Evolution of Depth Predictions during training of W-Net Connected Architecture.

6.2 Observations and Discussion

During the training process, several “phantom features” were observed in the models predicted output. For instance, the models were observed to produce silhouettes of chairs and tables in depth predictions of RGB images that did not include these items. This phenomenon exhibits learned features from previous training sessions indicating an underfitting effect which could be mitigated with proper data augmentation techniques. The “ghost chairs problem”, as it has become colloquially known, is addressed with tenuous hyperparameter tuning, longer training sessions, and more data.

Throughout the training process, depth prediction images were logged after each training batch. As expected, the images gradually transitioned from being generally amorphous to a well-defined image matching the edges in the RGB input image. In intermediate steps, irregular shapes can be seen forming over some of the distinct features in the input image as shown in Figures 9 and 10.

In some cases, the networks will produce outputs which exhibit fairly accurate segmentation results, but not entirely correct depth prediction. This can be observed in Figure 14. The trees in the background of the image below are mistakenly estimated to be closer to the observer than the two individuals in the picture.



Figure 11: Left to Right: RGB Input Images (Newly Seen, City), U-Net Depth Predictions (20 Epochs), W-Net Connected Depth Predictions (20 Epochs).




Figure 13: Left to Right: RGB Input Images (Newly Seen, City), W-Net Depth Predictions (20 Epochs), W-Net Connected Depth Predictions (20 Epochs).



Figure 12: Left to Right: RGB Input Images (Newly Seen, Vehicles), U-Net Depth Predictions (20 Epochs), W-Net Connected Depth Predictions (20 Epochs).




Figure 14: Correct Segmentation with Inaccurate Depth.

7 CONCLUSIONS

In conclusion, the best performing model in terms of MSE is W-Net Connected. Despite the closely clustered MSE values reported in Table 1, there is a stark difference in the individual model prediction accuracy when examined qualitatively. W-Net Connected appeared to outperform W-Net due to its connection to the original RGB input as expected from the literature review discussion in section 1.2.

W-Net Connected also provided the best depth predictions when given newly seen test images. U-Net reached its lowest validation loss in epoch 4, W-Net in epoch 16 and W-Net Connected in epoch 20.

Within W-Net Connected, the segmentation capabilities provided by a pre-trained component U-Net certainly appears to increase the degree of detail observed in the final prediction image. However, it is uncertain whether this is due to the improved segmentation or by simply using a larger network. W-Net and W-Net Connected logged the longest inference times with an average runtime of roughly 133ms. U-Net inference times were roughly 75% of that figure.

7.1 Future Work

To make pragmatic use of depth estimating neural networks, we intend to develop and maintain a Robot Operating System (ROS) package that can utilize state-of-the-art depth estimation networks to take in a live video feed and publish point-cloud depth topics for ODOA/3D mapping. This could simplify and democratize the integration of depth estimation neural networks into robotic systems. Rigorous characterization of the presented architectures will be completed with the full NYUD v2 datasets using losses as

defined in (Wang et al., 2015).

Features that can be added to the system that will improve its efficacy include Data Augmentation, Custom Loss Function Design, and further hyperparameter tuning. Quick improvements can be made by experimenting with more segmentation techniques and even further improved by incorporating semantic segmentation systems such as YOLOv3 or other FRCNN based networks as proven in (Jiao et al., 2018) and (Wang et al., 2015).

Further architectures to be tested and evaluated include the addition of further LSTM/GRU hidden units such as W-Net Connected + LSTM. Preliminary research in this paper and in accompanying references suggest the time series nature of moving depth image and the interpolated data points in the current datasets can benefit from memory units when deducing depth among sparsely populated depth maps. Another path to take, illuminated by the work done in this paper, is exploring the use of autoencoders for representation learning of depth data to improve the inference time of this system.

Finally, a review of appropriate loss functions will be conducted. While MSE is a standard and staple of measuring the success of depth-estimation, it is evident that the W-Net Connected model produces more coherent results than U-Net, yet scored lower during training and evaluation. From this result, we can look on to utilizing scoring functions such as Scale Invariant Loss, MSLE (Mean-Squared-Log-Error), and possibly custom loss functions that take into account more than relative or absolute difference between ground truth and predicted images.

The end result of the improvements above will be the practical real-time production of depth data fed into a generic package for autonomous robotic systems equipped with obstacle detection and avoidance.

ACKNOWLEDGEMENTS

We thank Worcester Polytechnic Institute (WPI) for providing computing resources and funding throughout this project.

REFERENCES

- Eigen, D., Puhrsch, C., and Fergus, R. (2014). Depth map prediction from a single image using a multi-scale deep network. In *Advances in neural information processing systems (NIPS)*.
- Farabet, C., Couprie, C., Najman, L., and LeCun, Y. (2012). Learning hierarchical features for scene label-

ing. *IEEE transactions on pattern analysis and machine intelligence*, 35(8):1915–1929.

- Geiger, A., Lenz, P., Stiller, C., and Urtasun, R. (2013). Vision meets robotics: The kitti dataset. In *The International Journal of Robotics Research*, vol. 32, no. 11, pp. 1231–1237.
- Godard, C., Aodha, O. M., and Brostow, G. J. (2017). Unsupervised monocular depth estimation with left-right consistency. In *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*.
- He, K., Zhang, X., Ren, S., and Sun, J. (2016). Deep residual learning for image recognition. In *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*.
- Jiao, J., Cao, Y., Song, Y., and Lau, R. (2018). Look deeper into depth: Monocular depth estimation with semantic booster and attention-driven loss. In *European Conference on Computer Vision*.
- Lee, J. and Kim, C. (2019). Monocular depth estimation using relative depth maps. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*.
- Lee, J. e. a. (2019). From big to small: Multi-scale local planar guidance for monocular depth estimation. In *arXiv preprint arXiv:1907.10326*.
- Liang, M. and Hu, X. (2015). Recurrent convolutional neural network for object recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 3367–3375.
- Ronneberger, O., Fischer, P., and Brox, T. (2015). U-net: Convolutional networks for biomedical image segmentation. In *International Conference on Medical image computing and computer-assisted intervention*. Springer.
- Silberman, N., Hoiem, D., Kohli, P., and Fergus, R. (2012). Indoor segmentation and support inference from rgb-d images. In *Computer Vision – ECCV 2012 Lecture Notes in Computer Science*, pp. 746–760.
- Wang, P., Shen, X., Lin, Z., Cohen, S., Price, B., and Yuille, A. (2015). Towards unified depth and semantic prediction from a single image. In *2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*.

APPENDIX

All network models discussed and developed in this research are available at: <https://github.com/mech0ctopus/depth-estimation>.